

# Point Estimation and Maximum Likelihood



CSE512 – Machine Learning, Spring 2018, Stony Brook University

Instructor: Minh Hoai Nguyen ([minhhoai@cs.stonybrook.edu](mailto:minhhoai@cs.stonybrook.edu))

Date: 24 Jan 2018

Many slides are by Carlos Guestrin at University of Washington

# Your First Consulting Job

- The company of a billionaire from Long Island is hiring. He receives too many applications that he cannot review them all. He consults you.
- You say: just flip a coin. A US quarter coin featuring the profile of George Washington. The text "UNITED STATES OF AMERICA" is at the top, "LIBERTY" is on the left, "IN GOD WE TRUST" is on the right, and "QUARTER DOLLAR" is at the bottom.
- He says: Not enough. A coin is 50/50 chance, so I still have to review 50% of the applications
- You say: flip a thumbtack instead. A 3D rendering of a thumbtack with a dark, oval-shaped head and a long, thin, gold-colored pin.
- He says: good idea. I have one right here.

# Your First Consulting Job

- He says: wait, if I flip it, what's the probability it will fall with the head up?
- You say: Please flip it a few times



- You say: The probability is:
- **He says: Why???**
- You say: Because...

# Thumbtack – Binomial Distribution

- $P(\text{Heads}) = \theta$ ,  $P(\text{Tails}) = 1 - \theta$
- Flips are i.i.d.:
  - Independent events
  - Identically distributed according to Binomial distribution
- Sequence  $\mathcal{D}$  of  $\alpha_H$  Heads and  $\alpha_T$  Tails

$$P(\mathcal{D} \mid \theta) = \theta^{\alpha_H} (1 - \theta)^{\alpha_T}$$

# Maximum Likelihood Estimation

- **Data:** Observed set  $D$  of  $\alpha_H$  Heads and  $\alpha_T$  Tails
- **Hypothesis:** Binomial distribution
- Learning  $\theta$  is an optimization problem
  - What's the objective function?
- MLE: Choose  $\theta$  that maximizes the probability of observed data:

$$\begin{aligned}\hat{\theta} &= \arg \max_{\theta} P(\mathcal{D} \mid \theta) \\ &= \arg \max_{\theta} \ln P(\mathcal{D} \mid \theta)\end{aligned}$$

# Your first learning algorithm

$$\begin{aligned}\hat{\theta} &= \arg \max_{\theta} \ln P(\mathcal{D} \mid \theta) \\ &= \arg \max_{\theta} \ln \theta^{\alpha_H} (1 - \theta)^{\alpha_T}\end{aligned}$$

- Set derivative to zero:  $\frac{d}{d\theta} \ln P(\mathcal{D} \mid \theta) = 0$

# How many flips do I need?

$$\hat{\theta}_{MLE} = \frac{\alpha_H}{\alpha_H + \alpha_T}$$

- Billionaire says: I flipped 3 heads and 2 tails.
- You say:  $\theta = 3/5$ , I can prove it!
- He says: What if I flipped 30 heads and 20 tails?
- You say: Same answer, I can prove it!
- **He says: What's better?**
- You say: Hmm... The more the merrier???
- He says: Is this why I am paying you the big bucks???

# Simple bound (based on Hoeffding's inequality)

- For  $N = \alpha_H + \alpha_T$ , and  $\hat{\theta}_{MLE} = \frac{\alpha_H}{\alpha_H + \alpha_T}$
- Let  $\theta^*$  be the true parameter, for any  $\epsilon > 0$ :

$$P(|\hat{\theta} - \theta^*| \geq \epsilon) \leq 2e^{-2N\epsilon^2}$$



# PAC Learning

- PAC: Probably Approximate Correct
- Billionaire says: I want to know the thumbtack parameter  $\theta$ , within  $\varepsilon = 0.1$ , with probability at least  $1 - \delta = 0.95$ . How many flips?

$$P(|\hat{\theta} - \theta^*| \geq \epsilon) \leq 2e^{-2N\epsilon^2}$$

# What about continuous variables?

- Billionaire says: If I want to predict the amount of fuel to fly from New York to Chicago, what can you do for me? It's a continuous variable!
- You say: Let me tell you about Gaussians...

$$P(x \mid \mu, \sigma) = \frac{1}{\sigma\sqrt{2\pi}} e^{\frac{-(x-\mu)^2}{2\sigma^2}}$$

# Gaussian Distribution Reviewed

Probability density function (pdf):  $P(x) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{1}{2\sigma^2}(x-\mu)^2}$

Notation:  $x \sim \mathcal{N}(\mu, \sigma^2)$

# Some Properties of Gaussians

## Affine transformation

$$x \sim \mathcal{N}(\mu, \sigma^2)$$

$$y = ax + b \Rightarrow y \sim \mathcal{N}(a\mu + b, a^2\sigma^2)$$

## Sum of Gaussians

$$y \sim \mathcal{N}(\mu_y, \sigma_y^2)$$

$$x \sim \mathcal{N}(\mu_x, \sigma_x^2)$$

$$z = x + y \Rightarrow z \sim \mathcal{N}(\mu_x + \mu_y, \sigma_x^2 + \sigma_y^2)$$

# Learning a Gaussian

- Given data samples  $D = \{x_1, \dots, x_n\}$
- Suppose the data comes from a Gaussian  $\mathcal{N}(\mu, \sigma^2)$
- Learn the parameters of the Gaussian

$$P(x|\mu, \sigma) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{1}{2\sigma^2}(x-\mu)^2}$$

# Maximum Likelihood Estimator (MLE)

The likelihood of data (probability)

$$P(D|\mu, \sigma) =$$

The log-likelihood

$$\log(P(D|\mu, \sigma)) =$$

# MLE for Mean

$$\frac{\partial \log(P(D|\mu, \sigma))}{\partial \mu} =$$

# MLE for Variance

$$\frac{\partial \log(P(D|\mu, \sigma))}{\partial \sigma} =$$



# MLE for Gaussian parameters

$$\hat{\mu}_{MLE} = \frac{1}{n} \sum_{i=1}^n x_i$$

$$\hat{\sigma}_{MLE} = \frac{1}{n} \sum_{i=1}^n (x_i - \hat{\mu})^2$$

# Biased and Unbiased Estimators

MLE of the mean is un-biased:

$$E_D[\hat{\mu}_{MLE}] = \mu$$

Expectation over all dataset  $D$  of  $n$  elements

Proof:

$$\begin{aligned} E_D[\hat{\mu}_{MLE}] &= E_D\left[\frac{1}{n} \sum_{i=1}^n x_i\right] \\ &= \frac{1}{n} \sum_{i=1}^n E[x_i] \\ &= \mu \end{aligned}$$

# Biased and Unbiased Estimators

MLE of the variance is biased:

$$E_D[\hat{\sigma}_{MLE}] \neq \sigma$$

Expectation over all dataset  $D$  of  $n$  elements

The unbiased estimator of variance

$$\hat{\sigma}_{unbiased} = \frac{1}{n-1} \sum_{i=1}^n (x_i - \hat{\mu})^2$$

Proof: exercise!

# Maximum A Posterior (MAP)

MLE:

$$\hat{\mu}, \hat{\sigma} = \operatorname{argmax}_{\mu, \sigma} P(D|\mu, \sigma)$$

Frequentist

MAP:

$$\hat{\mu}, \hat{\sigma} = \operatorname{argmax}_{\mu, \sigma} P(\mu, \sigma|D)$$

Bayesian

# MAP

$$P(\mu, \sigma | D) = \frac{P(D | \mu, \sigma) P(\mu, \sigma)}{P(D)}$$
$$\propto P(D | \mu, \sigma) P(\mu, \sigma)$$

MLE = MAP if we assume uniform prior for  $(\mu, \sigma)$

# MAP for Gaussian

Use conjugate priors for the parameters

- Mean: Gaussian prior
- Variance: Wishart Distribution

Prior for Mean

$$P(\mu) = \mathcal{N}(\mu|\mu_0, \lambda^2)$$

# MAP for Mean of Gaussian

$$\log(P(D|\mu, \sigma)P(\mu)) = \log(P(D|\mu, \sigma)) + \log(P(\mu))$$


To find MAP estimate, take derivative and set it to 0

# Prior for Thumbtack Problem


Data likelihood:  $P(\mathcal{D} \mid \theta) = \theta^{\alpha_H} (1 - \theta)^{\alpha_T}$

If we use Beta prior for the parameter  $P(\theta) = \frac{\theta^{\beta_H - 1} (1 - \theta)^{\beta_T - 1}}{B(\beta_H, \beta_T)}$


Beta(1,1)




Beta(2,2)



Beta(3,2)



Beta(30,20)





# MAP estimate for Thumbtack Problem

Data likelihood:  $P(\mathcal{D} \mid \theta) = \theta^{\alpha_H} (1 - \theta)^{\alpha_T}$

If we use Beta prior for the parameter  $P(\theta) = \frac{\theta^{\beta_H - 1} (1 - \theta)^{\beta_T - 1}}{B(\beta_H, \beta_T)}$

# What is Machine Learning? – Revisited

- Machine Learning is...
  - Collect some data
    - E.g., thumbtack flips, fuel consumption
  - Choose a hypothesis class or model
    - E.g., binomial, Gaussian
  - Choose a loss function
    - E.g., data likelihood, parameter likelihood
  - Choose an optimization procedure
    - E.g., set derivative to zero to obtain MLE

# Lecture Summary

- What Machine Learning is
- Maximum Likelihood Estimator (MLE)
- Maximum A Posterior Estimator (MAP)
- Binomial Distribution
- Gaussian Distribution